

## ✓ Lab 02: Basic NLP Preprocessing Techniques

**Course:** ITAI 2373 - Natural Language Processing

**Module:** 02 - Text Preprocessing

**Duration:** 2-3 hours

**Student Name:** \_\_\_\_Gregory Livingston\_\_\_\_

**Date:** \_\_\_\_6-10-2025\_\_\_\_

---

### Learning Objectives

By completing this lab, you will:

1. Understand the critical role of preprocessing in NLP pipelines
2. Master fundamental text preprocessing techniques
3. Compare different libraries and their approaches
4. Analyze the effects of preprocessing on text data
5. Build a complete preprocessing pipeline
6. Load and work with different types of text datasets

## ✓ Introduction to NLP Preprocessing

Natural Language Processing (NLP) preprocessing refers to the initial steps taken to clean and transform raw text data into a format that's more suitable for analysis by machine learning algorithms.

Why is preprocessing crucial?

1. **Standardization:** Ensures consistent text format across your dataset
2. **Noise Reduction:** Removes irrelevant information that could confuse algorithms
3. **Complexity Reduction:** Simplifies text to focus on meaningful patterns
4. **Performance Enhancement:** Improves the efficiency and accuracy of downstream tasks

Real-world Impact

Consider searching for "running shoes" vs "Running Shoes!" - without preprocessing, these might be treated as completely different queries. Preprocessing ensures they're recognized as equivalent.

### Conceptual Question 1

**Before we start coding, think about your daily interactions with text processing systems (search engines, chatbots, translation apps). What challenges do you think these systems face when processing human language? List at least 3 specific challenges and explain why each is problematic.**

*Double-click this cell to write your answer:*

**Challenge 1:** Like "bat" can be an animal or something to hit a ball. That can confuse the computer.

**Challenge 2:** We make typos or write "u" instead of "you," and the system might not know what we mean.


**Challenge 3:** When we say "that's lit", the computer might not know we mean "that's cool."

---

## ✓ Part 1: Environment Setup

We'll be working with two major NLP libraries:

- **NLTK (Natural Language Toolkit):** Comprehensive NLP library with extensive resources
- **spaCy:** Industrial-strength NLP with pre-trained models

 **Note:** Installation might take 2-3 minutes to complete.

```
# Step 1: Install Required Libraries
print("🔧 Installing NLP libraries...")
```

```
!pip install -q nltk spacy
!python -m spacy download en_core_web_sm
```

```
print("✅ Installation complete!")
```

```
🔄 Installing NLP libraries...
Collecting en-core-web-sm==3.8.0
  Downloading https://github.com/explosion/spacy-models/releases/download/en_core_web_sm-3.8.0/en_core_web_sm-3.8.0-py3-none-any.whl (12.8/12.8 MB 36.2 MB/s eta 0:00:00)
    Download and installation successful
You can now load the package via spacy.load('en_core_web_sm')
⚠️ Restart to reload dependencies
If you are in a Jupyter or Colab notebook, you may need to restart Python in order to load all the package's dependencies. You can do this by selecting the 'Restart kernel' or 'Restart runtime' option.
✅ Installation complete!
```

## 🧠 Conceptual Question 2

**Why do you think we need to install a separate language model (en\_core\_web\_sm) for spaCy? What components might this model contain that help with text processing? Think about what information a computer needs to understand English text.** We need to install the en\_core\_web\_sm model because it teaches the computer how to understand English. This model has parts that help it know what words mean, how they're used, and how they connect. It helps the computer tell the difference between names, places, actions, and feelings—kind of like giving it English-reading superpowers.

*Double-click this cell to write your answer:*

```
# Step 2: Import Libraries and Download NLTK Data
import nltk
import spacy
import string
import re
from collections import Counter

# Download essential NLTK data
print("📦 Downloading NLTK data packages...")
nltk.download('punkt')      # For tokenization
nltk.download('stopwords')  # For stop word removal
nltk.download('wordnet')    # For lemmatization
nltk.download('averaged_perceptron_tagger') # For POS tagging

print("\n✅ All imports and downloads completed!")
```

```
🔄 📦 Downloading NLTK data packages...
✅ All imports and downloads completed!
[nltk_data] Downloading package punkt to /root/nltk_data...
[nltk_data] Package punkt is already up-to-date!
[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data] Package stopwords is already up-to-date!
[nltk_data] Downloading package wordnet to /root/nltk_data...
[nltk_data] Package wordnet is already up-to-date!
[nltk_data] Downloading package averaged_perceptron_tagger to
[nltk_data] /root/nltk_data...
[nltk_data] Package averaged_perceptron_tagger is already up-to-
[nltk_data] date!
```

## 📁 Part 2: Sample Text Data

We'll work with different types of text to understand how preprocessing affects various text styles:

- Simple text
- Academic text (with citations, URLs)
- Social media text (with emojis, hashtags)
- News text (formal writing)
- Product reviews (informal, ratings)

```
# Step 3: Load Sample Texts
simple_text = "Natural Language Processing is a fascinating field of AI. It's amazing!"

academic_text = ""
```

```

Dr. Smith's research on machine-learning algorithms is groundbreaking!
She published 3 papers in 2023, focusing on deep neural networks (DNNs).
The results were amazing - accuracy improved by 15.7%!
"This is revolutionary," said Prof. Johnson.
Visit https://example.com for more info. #NLP #AI @university
"""

social_text = "OMG! Just tried the new coffee shop 🍷 SO GOOD!!! Highly recommend 👍 #coffee #yum 🥰"

news_text = """
The stock market experienced significant volatility today, with tech stocks leading the decline.
Apple Inc. (AAPL) dropped 3.2%, while Microsoft Corp. fell 2.8%.
"We're seeing a rotation out of growth stocks," said analyst Jane Doe from XYZ Capital.
"""

review_text = """
This laptop is absolutely fantastic! I've been using it for 6 months and it's still super fast.
The battery life is incredible - lasts 8-10 hours easily.
Only complaint: the keyboard could be better. Overall rating: 4.5/5 stars.
"""

# Store all texts
sample_texts = {
    "Simple": simple_text,
    "Academic": academic_text.strip(),
    "Social Media": social_text,
    "News": news_text.strip(),
    "Product Review": review_text.strip()
}

print("📄 Sample texts loaded successfully!")
for name, text in sample_texts.items():
    preview = text[:80] + "..." if len(text) > 80 else text
    print(f"📄 {name}: {preview}")

```



```

📄 Sample texts loaded successfully!

📄 Simple: Natural Language Processing is a fascinating field of AI. It's amazing!

📄 Academic: Dr. Smith's research on machine-learning algorithms is groundbreaking!
She publi...

📄 Social Media: OMG! Just tried the new coffee shop 🍷 SO GOOD!!! Highly recommend 👍 #coffee #yu...

📄 News: The stock market experienced significant volatility today, with tech stocks lead...

📄 Product Review: This laptop is absolutely fantastic! I've been using it for 6 months and it's st...

```

### 🤔 Conceptual Question 3

Looking at the different text types we've loaded, what preprocessing challenges do you anticipate for each type? For each text type below, identify at least 2 specific preprocessing challenges and explain why they might be problematic for NLP analysis.

*Double-click this cell to write your answer:*

**Simple text challenges:** 1. Too short, not enough words to understand. 2. Too easy – doesn't teach the computer much.

**Academic text challenges:** 1. Big words – hard for the computer to know them. 2. Long sentences – confusing to break apart.

**Social media text challenges:** 1. Emojis and hashtags – computer might not get them. 2. Slang – not normal spelling or grammar.

**News text challenges:** 1. Names and places – need to keep them special. 2. Formal words – hard to tell the feeling.

**Product review challenges:** 1. Opinions – hard to know if it's good or bad. 2. Mixed thoughts – one sentence can say good and bad.

## 📄 Part 3: Tokenization

### What is Tokenization?

Tokenization is the process of breaking down text into smaller, meaningful units called **tokens**. These tokens are typically words, but can also be sentences, characters, or subwords.

### Why is it Important?

- Most NLP algorithms work with individual tokens, not entire texts
- It's the foundation for all subsequent preprocessing steps
- Different tokenization strategies can significantly impact results

### Common Challenges:

- **Contractions:** "don't" → "do" + "n't" or "don't"?
- **Punctuation:** Keep with words or separate?
- **Special characters:** How to handle @, #, URLs?

```
import nltk
nltk.download('punkt')
```

```
[nltk_data] Downloading package punkt to /root/nltk_data...
[nltk_data]   Package punkt is already up-to-date!
True
```

```
# Step 4: Tokenization with NLTK
from nltk.tokenize import word_tokenize, sent_tokenize
import nltk
nltk.download('punkt_tab') # Download the missing resource
```

```
# Test on simple text
print("🔍 NLTK Tokenization Results")
print("=" * 40)
print(f"Original: {simple_text}")
```

```
# Word tokenization
nltk_tokens = word_tokenize(simple_text)
print(f"\nWord tokens: {nltk_tokens}")
print(f"Number of tokens: {len(nltk_tokens)}")
```

```
# Sentence tokenization
sentences = sent_tokenize(simple_text)
print(f"\nSentences: {sentences}")
print(f"Number of sentences: {len(sentences)}")
```

```
[nltk_data] Downloading package punkt_tab to /root/nltk_data...
[nltk_data]   Unzipping tokenizers/punkt_tab.zip.
🔍 NLTK Tokenization Results
=====
Original: Natural Language Processing is a fascinating field of AI. It's amazing!

Word tokens: ['Natural', 'Language', 'Processing', 'is', 'a', 'fascinating', 'field', 'of', 'AI', '.', 'It', 's', 'amazing', '!']
Number of tokens: 14

Sentences: ['Natural Language Processing is a fascinating field of AI.', "It's amazing!"]
Number of sentences: 2
```

Start coding or [generate](#) with AI.

### 📌 Conceptual Question 4

Examine the NLTK tokenization results above. How did NLTK handle the contraction "It's"? What happened to the punctuation marks? Do you think this approach is appropriate for all NLP tasks? Explain your reasoning.

*Double-click this cell to write your answer:*

**How "It's" was handled:** NLTK split "It's" into two parts: "It" and "s".

**Punctuation treatment:** It took away most punctuation and made them into separate tokens.

**Appropriateness for different tasks:** This is good for some tasks like counting words, but not good if you need the full meaning, like in a chatbot or translator.

```
# Step 5: Tokenization with spaCy
nlp = spacy.load('en_core_web_sm')

print("🔍 spaCy Tokenization Results")
print("=" * 40)
print(f"Original: {simple_text}")
```

```
# Process with spaCy
doc = nlp(simple_text)

# Extract tokens
spacy_tokens = [token.text for token in doc]
print(f"\nWord tokens: {spacy_tokens}")
print(f"Number of tokens: {len(spacy_tokens)}")

# Show detailed token information
print(f"\n🔍 Detailed Token Analysis:")
print(f"{'Token':<12} {'POS':<8} {'Lemma':<12} {'Is Alpha':<8} {'Is Stop':<8}")
print("-" * 50)
for token in doc:
    print(f"{'token.text':<12} {'token.pos_':<8} {'token.lemma_':<12} {'token.is_alpha':<8} {'token.is_stop':<8}")
```

🔗🔍 spaCy Tokenization Results

=====

Original: Natural Language Processing is a fascinating field of AI. It's amazing!

Word tokens: ['Natural', 'Language', 'Processing', 'is', 'a', 'fascinating', 'field', 'of', 'AI', '.', 'It', 's', 'amazing', '!']

Number of tokens: 14

🔍 Detailed Token Analysis:

Token	POS	Lemma	Is Alpha	Is Stop
Natural	PROPN	Natural	1	0
Language	PROPN	Language	1	0
Processing	NOUN	processing	1	0
is	AUX	be	1	1
a	DET	a	1	1
fascinating	ADJ	fascinating	1	0
field	NOUN	field	1	0
of	ADP	of	1	1
AI	PROPN	AI	1	0
.	PUNCT	.	0	0
It	PRON	it	1	1
's	AUX	be	0	1
amazing	ADJ	amazing	1	0
!	PUNCT	!	0	0

📌 Conceptual Question 5

Compare the NLTK and spaCy tokenization results. What differences do you notice? Which approach do you think would be better for different NLP tasks? Consider specific examples like sentiment analysis vs. information extraction.

Double-click this cell to write your answer:

- Key differences observed:**spaCy kept words together better and handled punctuation smarter than NLTK.
- Better for sentiment analysis:**NLTK is okay because it splits words and shows emotions clearly.
- Better for information extraction:**spaCy is better because it understands names and places more accurately.
- Overall assessment:**spaCy is smarter for hard tasks, but NLTK is good for learning and simple stuff.

```
# Step 6: Test Tokenization on Complex Text
print("🔗 Testing on Social Media Text")
print("-" * 40)
print(f"Original: {social_text}")

# NLTK approach
social_nltk_tokens = word_tokenize(social_text)
print(f"\nNLTK tokens: {social_nltk_tokens}")

# spaCy approach
social_doc = nlp(social_text)
social_spacy_tokens = [token.text for token in social_doc]
print(f"spaCy tokens: {social_spacy_tokens}")

print(f"\n🔍 Comparison:")
print(f"NLTK token count: {len(social_nltk_tokens)}")
print(f"spaCy token count: {len(social_spacy_tokens)}")
```

📌 Conceptual Question 6

Looking at how the libraries handled social media text (emojis, hashtags), which library seems more robust for handling "messy" real-world text? What specific advantages do you notice? How might this impact a real-world application like social media sentiment analysis?

Double-click this cell to write your answer:

**More robust library:**spaCy

**Specific advantages:**spaCy understands emojis and hashtags better and keeps things organized.

**Impact on sentiment analysis:**It helps find feelings in tweets or posts more correctly because it sees the whole meaning, not just pieces.

---

## 📌 Part 4: Stop Words Removal

### What are Stop Words?

Stop words are common words that appear frequently in a language but typically don't carry much meaningful information about the content. Examples include "the", "is", "at", "which", "on", etc.

### Why Remove Stop Words?

1. **Reduce noise** in the data
2. **Improve efficiency** by reducing vocabulary size
3. **Focus on content words** that carry semantic meaning

### When NOT to Remove Stop Words?

- **Sentiment analysis:** "not good" vs "good" - the "not" is crucial!
- **Question answering:** "What is the capital?" - "what" and "is" provide context

```
# Step 7: Explore Stop Words Lists
from nltk.corpus import stopwords

# Get NLTK English stop words
nltk_stopwords = set(stopwords.words('english'))
print(f"📌 NLTK has {len(nltk_stopwords)} English stop words")
print(f"First 20: {sorted(list(nltk_stopwords))[:20]}")

# Get spaCy stop words
spacy_stopwords = nlp.Defaults.stop_words
print(f"\n📌 spaCy has {len(spacy_stopwords)} English stop words")
print(f"First 20: {sorted(list(spacy_stopwords))[:20]}")

# Compare the lists
common_stopwords = nltk_stopwords.intersection(spacy_stopwords)
nltk_only = nltk_stopwords - spacy_stopwords
spacy_only = spacy_stopwords - nltk_stopwords

print(f"\n🔍 Comparison:")
print(f"Common stop words: {len(common_stopwords)}")
print(f"Only in NLTK: {len(nltk_only)} - Examples: {sorted(list(nltk_only))[:5]}")
print(f"Only in spaCy: {len(spacy_only)} - Examples: {sorted(list(spacy_only))[:5]}")

🔄 📌 NLTK has 198 English stop words
First 20: ['a', 'about', 'above', 'after', 'again', 'against', 'ain', 'all', 'am', 'an', 'and', 'any', 'are', 'aren', "aren't", 'as', 'a

📌 spaCy has 326 English stop words
First 20: ["'d", "'ll", "'m", "'re", "'s", "'ve", 'a', 'about', 'above', 'across', 'after', 'afterwards', 'again', 'against', 'all', 'al

🔍 Comparison:
Common stop words: 123
Only in NLTK: 75 - Examples: ['ain', 'aren', "aren't", 'couldn', "couldn't"]
Only in spaCy: 203 - Examples: ["'d", "'ll", "'m", "'re", "'s"]
```

---

## 🧠 Conceptual Question 7

Why do you think NLTK and spaCy have different stop word lists? Look at the examples of words that are only in one list - do you agree with these choices? Can you think of scenarios where these differences might significantly impact your NLP results?

Double-click this cell to write your answer:

**Reasons for differences:**They were made by different teams and used for different things, so they picked different words.

**Agreement with choices:** Yeah, it makes sense. spaCy has more words like “I’m” and “you’re” which are common online.

**Scenarios where differences matter:** If you’re working on tweets or texts, spaCy might do better. But for school books or essays, NLTK could be fine.

```
# Step 8: Remove Stop Words with NLTK
# Test on simple text
original_tokens = nltk_tokens # From earlier tokenization
filtered_tokens = [word for word in original_tokens if word.lower() not in nltk_stopwords]
```

```
print("🧪 NLTK Stop Word Removal")
print("=" * 40)
print(f"Original: {simple_text}")
print(f"\nOriginal tokens ({len(original_tokens)}): {original_tokens}")
print(f"After removing stop words ({len(filtered_tokens)}): {filtered_tokens}")
```

```
# Show which words were removed
removed_words = [word for word in original_tokens if word.lower() in nltk_stopwords]
print(f"\nRemoved words: {removed_words}")
```

```
# Calculate reduction percentage
reduction = (len(original_tokens) - len(filtered_tokens)) / len(original_tokens) * 100
print(f"Vocabulary reduction: {reduction:.1f}%")
```

```
🔄 🧪 NLTK Stop Word Removal
=====
Original: Natural Language Processing is a fascinating field of AI. It's amazing!

Original tokens (14): ['Natural', 'Language', 'Processing', 'is', 'a', 'fascinating', 'field', 'of', 'AI', '.', 'It', "'s", 'amazing', '']
After removing stop words (10): ['Natural', 'Language', 'Processing', 'fascinating', 'field', 'AI', '.', "'s", 'amazing', '!']

Removed words: ['is', 'a', 'of', 'It']
Vocabulary reduction: 28.6%
```

```
# Step 9: Remove Stop Words with spaCy
doc = nlp(simple_text)
spacy_filtered = [token.text for token in doc if not token.is_stop and not token.is_punct]
```

```
print("🧪 spaCy Stop Word Removal")
print("=" * 40)
print(f"Original: {simple_text}")
print(f"\nOriginal tokens ({len(spacy_tokens)}): {spacy_tokens}")
print(f"After removing stop words & punctuation ({len(spacy_filtered)}): {spacy_filtered}")
```

```
# Show which words were removed
spacy_removed = [token.text for token in doc if token.is_stop or token.is_punct]
print(f"\nRemoved words: {spacy_removed}")
```

```
# Calculate reduction percentage
spacy_reduction = (len(spacy_tokens) - len(spacy_filtered)) / len(spacy_tokens) * 100
print(f"Vocabulary reduction: {spacy_reduction:.1f}%")
```

```
🔄 🧪 spaCy Stop Word Removal
=====
Original: Natural Language Processing is a fascinating field of AI. It's amazing!

Original tokens (14): ['Natural', 'Language', 'Processing', 'is', 'a', 'fascinating', 'field', 'of', 'AI', '.', 'It', "'s", 'amazing', '']
After removing stop words & punctuation (7): ['Natural', 'Language', 'Processing', 'fascinating', 'field', 'AI', 'amazing']

Removed words: ['is', 'a', 'of', '.', 'It', "'s", '!']
Vocabulary reduction: 50.0%
```

## 🤖 Conceptual Question 8

**Compare the NLTK and spaCy stop word removal results. Which approach removed more words? Do you think removing punctuation (as spaCy did) is always a good idea? Give a specific example where keeping punctuation might be important for NLP analysis.**

*Double-click this cell to write your answer:*

**Which removed more:** spaCy took out more words than NLTK.

**Punctuation removal assessment:** Taking out punctuation is not always good.

**Example where punctuation matters:** In "Wait... what?!" the dots and question mark show surprise and confusion. That helps tell the mood.

## ✓ 🌱 Part 5: Lemmatization and Stemming

### What is Lemmatization?

Lemmatization reduces words to their base or dictionary form (called a **lemma**). It considers context and part of speech to ensure the result is a valid word.

### What is Stemming?

Stemming reduces words to their root form by removing suffixes. It's faster but less accurate than lemmatization.

### Key Differences:

Aspect	Stemming	Lemmatization
Speed	Fast	Slower
Accuracy	Lower	Higher
Output	May be non-words	Always valid words
Context	Ignores context	Considers context

### Examples:

- **"running"** → Stem: "run", Lemma: "run"
- **"better"** → Stem: "better", Lemma: "good"
- **"was"** → Stem: "wa", Lemma: "be"

```
# Step 10: Stemming with NLTK
from nltk.stem import PorterStemmer

stemmer = PorterStemmer()

# Test words that demonstrate stemming challenges
test_words = ['running', 'runs', 'ran', 'better', 'good', 'best', 'flying', 'flies', 'was', 'were', 'cats', 'dogs']

print("🌱 Stemming Demonstration")
print("=" * 30)
print(f"{'Original':<12} {'Stemmed':<12}")
print("-" * 25)

for word in test_words:
    stemmed = stemmer.stem(word)
    print(f"{word:<12} {stemmed:<12}")

# Apply to our sample text
sample_tokens = [token for token in nltk_tokens if token.isalpha()]
stemmed_tokens = [stemmer.stem(token.lower()) for token in sample_tokens]

print(f"\n🌱 Applied to sample text:")
print(f"Original: {sample_tokens}")
print(f"Stemmed: {stemmed_tokens}")

🔄 🌱 Stemming Demonstration
=====
Original      Stemmed
-----
running      run
runs         run
ran          ran
better       better
good         good
best        best
flying      fli
flies       fli
was         wa
were        were
cats        cat
dogs        dog

🌱 Applied to sample text:
Original: ['Natural', 'Language', 'Processing', 'is', 'a', 'fascinating', 'field', 'of', 'AI', 'It', 'amazing']
Stemmed: ['natur', 'languag', 'process', 'is', 'a', 'fascin', 'field', 'of', 'ai', 'it', 'amaz']
```



## ✓ 🤖 Conceptual Question 9

Look at the stemming results above. Can you identify any cases where stemming produced questionable results? For example, how were "better" and "good" handled? Do you think this is problematic for NLP applications? Explain your reasoning.

Double-click this cell to write your answer:

**Questionable results identified:** Yes, words like "better" and "good" didn't get changed the right way.

**Assessment of "better" and "good":** Stemming didn't understand that "better" and "good" mean something similar.

**Impact on NLP applications:** This can confuse the computer and make it think the words are different when they really mean the same thing.

```
# Step 11: Lemmatization with spaCy
print("🦋 spaCy Lemmatization Demonstration")
print("=" * 40)

# Test on a complex sentence
complex_sentence = "The researchers were studying the effects of running and swimming on better performance."
doc = nlp(complex_sentence)

print(f"Original: {complex_sentence}")
print(f"\n{'Token':<15} {'Lemma':<15} {'POS':<10} {'Explanation':<20}")
print("-" * 65)

for token in doc:
    if token.is_alpha:
        explanation = "No change" if token.text.lower() == token.lemma_ else "Lemmatized"
        print(f"{token.text:<15} {token.lemma_:<15} {token.pos_:<10} {explanation:<20}")

# Extract lemmas
lemmas = [token.lemma_.lower() for token in doc if token.is_alpha and not token.is_stop]
print(f"\n📋 Lemmatized tokens (no stop words): {lemmas}")
```

```
🦋 🦋 spaCy Lemmatization Demonstration
=====
Original: The researchers were studying the effects of running and swimming on better performance.
```

Token	Lemma	POS	Explanation
The	the	DET	No change
researchers	researcher	NOUN	Lemmatized
were	be	AUX	Lemmatized
studying	study	VERB	Lemmatized
the	the	DET	No change
effects	effect	NOUN	Lemmatized
of	of	ADP	No change
running	run	VERB	Lemmatized
and	and	CCONJ	No change
swimming	swim	VERB	Lemmatized
on	on	ADP	No change
better	well	ADJ	Lemmatized
performance	performance	NOUN	No change

```
📋 Lemmatized tokens (no stop words): ['researcher', 'study', 'effect', 'run', 'swim', 'well', 'performance']
```

```
# Step 12: Compare Stemming vs Lemmatization
comparison_words = ['better', 'running', 'studies', 'was', 'children', 'feet']

print("🦋 Stemming vs Lemmatization Comparison")
print("=" * 50)
print(f"{'Original':<12} {'Stemmed':<12} {'Lemmatized':<12}")
print("-" * 40)

for word in comparison_words:
    # Stemming
    stemmed = stemmer.stem(word)

    # Lemmatization with spaCy
    doc = nlp(word)
    lemmatized = doc[0].lemma_

    print(f"{word:<12} {stemmed:<12} {lemmatized:<12}")
```

```
🦋 🦋 Stemming vs Lemmatization Comparison
=====
Original      Stemmed      Lemmatized
```

better	better	well
running	run	run
studies	studi	study
was	wa	be
children	children	child
feet	feet	foot

## 🤖 Conceptual Question 10

Compare the stemming and lemmatization results. Which approach do you think is more suitable for:

1. **A search engine** (where speed is crucial and you need to match variations of words)?Stemming
2. **A sentiment analysis system** (where accuracy and meaning preservation are important)?Lemmatization
3. **A real-time chatbot** (where both speed and accuracy matter)?Lemmatization

Explain your reasoning for each choice.

Double-click this cell to write your answer:

1. **Search engine:**Stemming is better because it's faster and just needs to match words like "run," "running," and "ran" quickly.
2. **Sentiment analysis:**Lemmatization is better because it keeps the real meaning of words like "good" and "better," which helps the system understand feelings better.
3. **Real-time chatbot:**Lemmatization is better, but if it needs to be super fast, maybe use a mix of both. It has to be quick and still make sense!

## ✓ ✂ Part 6: Text Cleaning and Normalization

### What is Text Cleaning?

Text cleaning involves removing or standardizing elements that might interfere with analysis:

- **Case normalization** (converting to lowercase)
- **Punctuation removal**
- **Number handling** (remove, replace, or normalize)
- **Special character handling** (URLs, emails, mentions)
- **Whitespace normalization**

### Why is it Important?

- Ensures consistency across your dataset
- Reduces vocabulary size
- Improves model performance
- Handles edge cases in real-world data

```
# Step 13: Basic Text Cleaning
def basic_clean_text(text):
    """Apply basic text cleaning operations"""
    # Convert to lowercase
    text = text.lower()

    # Remove extra whitespace
    text = re.sub(r'\s+', ' ', text).strip()

    # Remove punctuation
    text = text.translate(str.maketrans('', '', string.punctuation))

    # Remove numbers
    text = re.sub(r'\d+', '', text)

    # Remove extra spaces again
    text = re.sub(r'\s+', ' ', text).strip()

    return text

# Test basic cleaning
test_text = " Hello WORLD!!! This has 123 numbers and extra spaces. "
cleaned = basic_clean_text(test_text)
```

```
print("🔪 Basic Text Cleaning")
print("=" * 30)
print(f"Original: '{test_text}'")
print(f"Cleaned: '{cleaned}'")
print(f"Length reduction: {(len(test_text) - len(cleaned))/len(test_text)*100:.1f}%")
```

#### 🔪 Basic Text Cleaning

```
=====
Original: ' Hello WORLD!!! This has 123 numbers and extra spaces. '
Cleaned: 'hello world this has numbers and extra spaces'
Length reduction: 26.2%
```

#### # Step 14: Advanced Cleaning for Social Media

```
def advanced_clean_text(text):
    """Apply advanced cleaning for social media and web text"""
    # Remove URLs
    text = re.sub(r'http\S+|www\S+|https\S+', '', text, flags=re.MULTILINE)

    # Remove email addresses
    text = re.sub(r'\S+@\S+', '', text)

    # Remove mentions (@username)
    text = re.sub(r'@\w+', '', text)

    # Convert hashtags (keep the word, remove #)
    text = re.sub(r'#(\w+)', r'\1', text)

    # Remove emojis (basic approach)
    emoji_pattern = re.compile("[
        u"\U0001F600-\U0001F64F" # emoticons
        u"\U0001F300-\U0001F5FF" # symbols & pictographs
        u"\U0001F680-\U0001F6FF" # transport & map symbols
        u"\U0001F1E0-\U0001F1FF" # flags
    "]" + re.UNICODE)
    text = emoji_pattern.sub(r'', text)

    # Convert to lowercase and normalize whitespace
    text = text.lower()
    text = re.sub(r'\s+', ' ', text).strip()

    return text

# Test on social media text
print("🔪 Advanced Cleaning on Social Media Text")
print("=" * 45)
print(f"Original: {social_text}")

cleaned_social = advanced_clean_text(social_text)
print(f"Cleaned: {cleaned_social}")
print(f"Length reduction: {(len(social_text) - len(cleaned_social))/len(social_text)*100:.1f}%")
```

#### 🔪 Advanced Cleaning on Social Media Text

```
=====
Original: OMG! Just tried the new coffee shop ☕ SO GOOD!!! Highly recommend 👍 #coffee #yum 🍷
Cleaned: omg! just tried the new coffee shop ☕ so good!!! highly recommend coffee yum 🍷
Length reduction: 7.2%
```

## 🤔 Conceptual Question 11

Look at the advanced cleaning results for the social media text. What information was lost during cleaning? Can you think of scenarios where removing emojis and hashtags might actually hurt your NLP application? What about scenarios where keeping them would be beneficial?

*Double-click this cell to write your answer:*

**Information lost:** The feelings showed was off

**Scenarios where removal hurts:** If you're trying to find out how someone feels about coffee, removing emojis like that might make it harder to know if they liked it

**Scenarios where keeping helps:** For apps that check emotions or reviews, keeping emojis can help the computer know someone is happy or excited

---

## ✓ Part 7: Building a Complete Preprocessing Pipeline

Now let's combine everything into a comprehensive preprocessing pipeline that you can customize based on your needs.

Pipeline Components:

1. **Text cleaning** (basic or advanced)
2. **Tokenization** (NLTK or spaCy)
3. **Stop word removal** (optional)
4. **Lemmatization/Stemming** (optional)
5. **Additional filtering** (length, etc.)

```
# Step 15: Complete Preprocessing Pipeline
def preprocess_text(text,
                    clean_level='basic',    # 'basic' or 'advanced'
                    remove_stopwords=True,
                    use_lemmatization=True,
                    use_stemming=False,
                    min_length=2):
    """
    Complete text preprocessing pipeline
    """
    # Step 1: Clean text
    if clean_level == 'basic':
        cleaned_text = basic_clean_text(text)
    else:
        cleaned_text = advanced_clean_text(text)

    # Step 2: Tokenize
    if use_lemmatization:
        # Use spaCy for lemmatization
        doc = nlp(cleaned_text)
        tokens = [token.lemma_.lower() for token in doc if token.is_alpha]
    else:
        # Use NLTK for basic tokenization
        tokens = word_tokenize(cleaned_text)
        tokens = [token for token in tokens if token.isalpha()]




    # Step 3: Remove stop words
    if remove_stopwords:
        if use_lemmatization:
            tokens = [token for token in tokens if token not in spacy_stopwords]
        else:
            tokens = [token.lower() for token in tokens if token.lower() not in nltk_stopwords]

    # Step 4: Apply stemming if requested
    if use_stemming and not use_lemmatization:
        tokens = [stemmer.stem(token.lower()) for token in tokens]

    # Step 5: Filter by length
    tokens = [token for token in tokens if len(token) >= min_length]

    return tokens

print("🔧 Preprocessing Pipeline Created!")
print("✅ Ready to test different configurations.")

  Preprocessing Pipeline Created!
 Ready to test different configurations.
```

```
# Step 16: Test Different Pipeline Configurations
test_text = sample_texts["Product Review"]
print(f"🌀 Testing on: {test_text[:100]}...")
print("=" * 60)

# Configuration 1: Minimal processing
minimal = preprocess_text(test_text,
                          clean_level='basic',
                          remove_stopwords=False,
                          use_lemmatization=False,
                          use_stemming=False)
print(f"\n1. Minimal processing ({len(minimal)} tokens):")
print(f"    {minimal[:10]}...")
```

```
# Configuration 2: Standard processing
standard = preprocess_text(test_text,
                           clean_level='basic',
                           remove_stopwords=True,
                           use_lemmatization=True)
print(f"\n2. Standard processing ({len(standard)} tokens):")
print(f"    {standard[:10]}...")

# Configuration 3: Aggressive processing
aggressive = preprocess_text(test_text,
                             clean_level='advanced',
                             remove_stopwords=True,
                             use_lemmatization=False,
                             use_stemming=True,
                             min_length=3)
print(f"\n3. Aggressive processing ({len(aggressive)} tokens):")
print(f"    {aggressive[:10]}...")

# Show reduction percentages
original_count = len(word_tokenize(test_text))
print(f"\n📊 Token Reduction Summary:")
print(f"    Original: {original_count} tokens")
print(f"    Minimal: {len(minimal)} ({(original_count-len(minimal))/original_count*100:.1f}% reduction)")
print(f"    Standard: {len(standard)} ({(original_count-len(standard))/original_count*100:.1f}% reduction)")
print(f"    Aggressive: {len(aggressive)} ({(original_count-len(aggressive))/original_count*100:.1f}% reduction)")

🔄 Testing on: This laptop is absolutely fantastic! I've been using it for 6 months and it's still super fast.
The ...
=====

1. Minimal processing (34 tokens):
   ['this', 'laptop', 'is', 'absolutely', 'fantastic', 'ive', 'been', 'using', 'it', 'for']...

2. Standard processing (18 tokens):
   ['laptop', 'absolutely', 'fantastic', 've', 'use', 'month', 'super', 'fast', 'battery', 'life']...

3. Aggressive processing (21 tokens):
   ['laptop', 'absolut', 'fantast', 'use', 'month', 'still', 'super', 'fast', 'batteri', 'life']...

📊 Token Reduction Summary:
Original: 47 tokens
Minimal: 34 (27.7% reduction)
Standard: 18 (61.7% reduction)
Aggressive: 21 (55.3% reduction)
```

## ✓ 🤖 Conceptual Question 12

Compare the three pipeline configurations (Minimal, Standard, Aggressive). For each configuration, analyze:

1. What information was preserved?
2. What information was lost?
3. What type of NLP task would this configuration be best suited for?

Double-click this cell to write your answer:

### Minimal Processing:

- Preserved: Most of the words, including small ones like "is" or "the"
- Lost: Just a few unimportant words
- Best for: When we want to keep full meaning, like in story writing or chatbots

### Standard Processing:

- Preserved: Only the main words that matter
- Lost: Small words and some extras like punctuation
- Best for: Tasks like finding out how someone feels (sentiment analysis)

### Aggressive Processing:

- Preserved: Just the strongest, most important words
- Lost: A lot of details and smaller ideas
- Best for: Fast searches or when we just need the big idea, like tagging topics

---

```
# Step 17: Comprehensive Analysis Across Text Types
print("📊 Comprehensive Preprocessing Analysis")
```

```

print("=" * 50)

# Test standard preprocessing on all text types
results = {}
for name, text in sample_texts.items():
    original_tokens = len(word_tokenize(text))
    processed_tokens = preprocess_text(text,
                                       clean_level='basic',
                                       remove_stopwords=True,
                                       use_lemmatization=True)

    reduction = (original_tokens - len(processed_tokens)) / original_tokens * 100
    results[name] = {
        'original': original_tokens,
        'processed': len(processed_tokens),
        'reduction': reduction,
        'sample': processed_tokens[:8]
    }

print(f"\n 📄 {name}:")
print(f"   Original: {original_tokens} tokens")
print(f"   Processed: {len(processed_tokens)} tokens ({reduction:.1f}% reduction)")
print(f"   Sample: {processed_tokens[:8]}")

# Summary table
print(f"\n\n 📊 Summary Table")
print(f"{'Text Type':<15} {'Original':<10} {'Processed':<10} {'Reduction':<10}")
print(f"-" * 50)
for name, data in results.items():
    print(f"{'name':<15} {'data['original']':<10} {'data['processed']':<10} {'data['reduction']':<10.1f}%")

```

## 🔗 Comprehensive Preprocessing Analysis

```

📄 Simple:
Original: 14 tokens
Processed: 7 tokens (50.0% reduction)
Sample: ['natural', 'language', 'processing', 'fascinating', 'field', 'ai', 'amazing']

📄 Academic:
Original: 61 tokens
Processed: 26 tokens (57.4% reduction)
Sample: ['dr', 'smith', 'research', 'machinelearning', 'algorithm', 'groundbreake', 'publish', 'paper']

📄 Social Media:
Original: 22 tokens
Processed: 10 tokens (54.5% reduction)
Sample: ['omg', 'try', 'new', 'coffee', 'shop', 'good', 'highly', 'recommend']

📄 News:
Original: 51 tokens
Processed: 25 tokens (51.0% reduction)
Sample: ['stock', 'market', 'experience', 'significant', 'volatility', 'today', 'tech', 'stock']

📄 Product Review:
Original: 47 tokens
Processed: 18 tokens (61.7% reduction)
Sample: ['laptop', 'absolutely', 'fantastic', 've', 'use', 'month', 'super', 'fast']

```

📊 Summary Table				
Text Type	Original	Processed	Reduction	
Simple	14	7	50.0	%
Academic	61	26	57.4	%
Social Media	22	10	54.5	%
News	51	25	51.0	%
Product Review	47	18	61.7	%

## 😬 Final Conceptual Question 13

Looking at the comprehensive analysis results across all text types:

1. **Which text type was most affected by preprocessing?** Why do you think this happened?
2. **Which text type was least affected?** What does this tell you about the nature of that text? Simple Text
3. **If you were building an NLP system to analyze customer reviews for a business, which preprocessing approach would you choose and why?**

#### 4. What are the main trade-offs you need to consider when choosing preprocessing techniques for any NLP project?

Double-click this cell to write your answer:

**1. Most affected text type:**Product reviews lost the most words because they had lots of extra or emotional words that weren't needed.

**2. Least affected text type:**Simple text didn't lose as much, because it already had fewer and more important words.

**3. For customer review analysis:**I'd use the standard cleaning style. It keeps the helpful words but removes the boring stuff, so we can still understand how people feel.

**4. Main trade-offs to consider:**If you clean too much, you might lose meaning. If you don't clean enough, the computer might get confused. You have to pick what's best for your task!

---

## Lab Summary and Reflection

Congratulations! You've completed a comprehensive exploration of NLP preprocessing techniques.

### Key Concepts You've Mastered:

1. **Text Preprocessing Fundamentals** - Understanding why preprocessing is crucial
2. **Tokenization Techniques** - NLTK vs spaCy approaches and their trade-offs
3. **Stop Word Management** - When to remove them and when to keep them
4. **Morphological Processing** - Stemming vs lemmatization for different use cases
5. **Text Cleaning Strategies** - Basic vs advanced cleaning for different text types
6. **Pipeline Design** - Building modular, configurable preprocessing systems

### Real-World Applications:

These techniques form the foundation for search engines, chatbots, sentiment analysis, document classification, machine translation, and information extraction systems.

### Key Insights to Remember:

- **No Universal Solution:** Different NLP tasks require different preprocessing approaches
  - **Trade-offs Are Everywhere:** Balance information preservation with noise reduction
  - **Context Matters:** The same technique can help or hurt depending on your use case
  - **Experimentation Is Key:** Always test and measure impact on your specific task
- 

## Excellent work completing Lab 02! 🎉

For your reflection journal, focus on the insights you gained about when and why to use different techniques, the challenges you encountered, and connections you made to real-world applications.